AWS Certified Solutions Architect: Associate - 5.0 ELB, CloudWatch and Auto Scaling

filename: amazon-acsaa-5-1-elb_and_auto_scaling
Title: ELB and Auto Scaling
Subtitle: AWS Certified Solutions Architect: Associate

5.1 ELB and Auto Scaling

- Elastic Load Balancing
  - Distributes traffic load across multiple instances
  - Managed by AWS
    - Highly available
    - Scales in/out automatically
  - Types
    - Internet facing
      - Uses public DNS/IP for routing
      - You should always use the DNS name
      - Publicly available
    - Internal load balancers
      - Used in VPCs with private subnets
    - HTTPS load balancers
      - SSL is terminated on the load balancer (SSL offload)
      - Certificate must be installed on the ELB
      - Does not support *Server Name Indication* (SNI)
        - Causes issues if multiple sites/domains are hosted behind the ELB
        - SSL cert must contain *Subject Alternative Name* (SAN) for each domain behind the ELB
  - Listeners
    - ELB only balances on specified methods
    - Supported methods
      - HTTP
      - HTTPS
      - TCP
      - SSL
  - Configuration
    - Idle connection timeout
      - 60 seconds by default
      - With HTTP/HTTPS the *keep-alive* option can be used
      - Otherwise, idle connections may reconnect to a different instance
    - Cross-zone Load Balancing
    - Connection Draining
      - Allows the ELB to stop sending connections to instances that are deregistering or unhealthy
      - Connections are held open for 300 seconds by default
      - Attempts to complete transaction if the instance becomes healthy again
    - Proxy protocol
      - Human readable header is attached to inbound traffic
      - Contains IP and port information from original traffic
      - Not usually necessary and can even cause issues
    - Sticky sessions
      - Also called *session affinity*
      - Normally, all connections are routed independantly
      - Sticky sessions route all requests from a client to the same instance
      - Uses a cookie to track the sessions
    - Health check
      - Testing method to determine if an instance is healthy
      - Methods
        - Ping
        - Connection
        - Web page
- Amazon CloudWatch
  - Monitoring service hosted inside of AWS
  - Provides tracking and alerting
  - Actions can be extended with AWS Simple Notification Service, Lambda, etc.
  - Used in conjunction with auto scaling
- Auto Scaling
  - Allows scaling instances in/out based on performance data gathered from CloudWatch
  - Auto Scaling Plans
    - Maintain current levels
      - Enforces a minimum # of instances

        - Manual scaling
        - Scheduled scaling
        - Dynamic scaling
    - Limits
        - `aws autoscaling describe-account-limits`
        - 100 launch configurations per region
        - 20 EC2 instances per region
    - Components (Diagram)
        - Launch configuration
            - Name
            - AMI
            - Instance Type
            - Security Group
            - Instance Key Pair
        - Auto Scaling Group
            - Name
            - Launch configuration
            - Availability Zones
            - Minimum Size
            - Desired Capacity
            - Maximum Capacity
            - Load Balancers
        - Scaling Policy (2 diagrams)
            - Associates CloudWatch alarms with scaling operations
    - Scale out quickly, scale in slowly